

# Web-Scraping

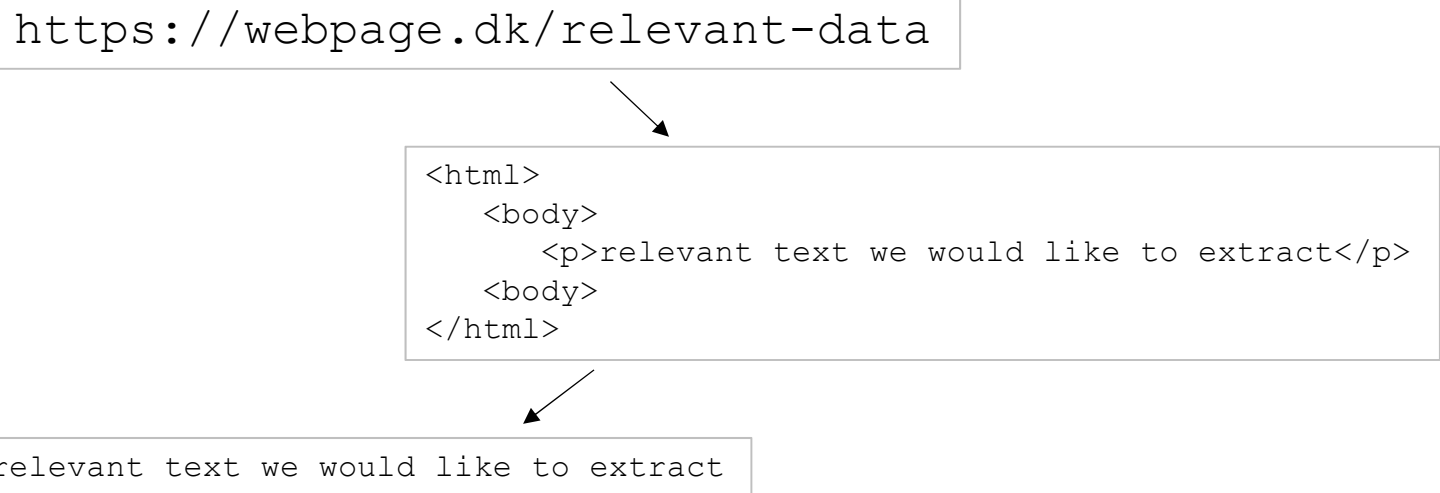
# Agenda

- Introduction
- Copyright
- Internet and WWW - brief history
- HTML
- CSS
- Python web-scraping
- Exercises

# What is Web-Scraping?

- Extract relevant text or images from a website
- Three steps
  1. Fetch the web-page (HTML-document)
  2. Parse the HTML
  3. Query and extract the relevant parts

`https://webpage.dk/relevant-data`




```
<html>
  <body>
    <p>relevant text we would like to extract</p>
  </body>
</html>
```

The diagram illustrates the web scraping process. It starts with a URL in a box at the top left. An arrow points from this box to a box containing a sample HTML document. Another arrow points from the HTML box to a box at the bottom left containing the extracted text. This visualizes the three steps: fetching the page, parsing the HTML, and extracting the relevant parts.

`relevant text we would like to extract`

# Copyright

- Even though publicly available many websites have copyright restrictions which must be followed:
  - look at Terms of Service
  - look at /robots.txt
- If we are allowed to scrape data, the following rules must still be followed:
  - **do not** re-distribute collected data without a clear written statement saying you are allowed to
  - **always** make clear where data comes from when used in papers etc. -> use quotation and references




# Copyright Examples

- Take a look at:  
<https://www.themoviedb.org/terms-of-use>
- <https://www.fandom.com/terms-of-use>
- <https://www.imdb.com/conditions>
- What are we allowed to do?

# The Internet and WWW

- 1960'es – Arpanet (several other networks followed)
- 1980'es – Different networks switching to the tcp/ip protocol -> global internet
- Protocols:
  - 1970'es – FTP (file transfer)
  - 1980'es – SMTP (email)
  - 1989 – HTTP (WWW)




# Tim Berners Lee and WWW

- CERN
  - Communication problems
- Solution WWW
  - HTML
    - Hypertext (Vannemar Bush, Ted Nelson etc.)
  - HTTP-protocol
  - Browser
  - Web-Server



# Client - Server

- For a website to become publicly available it must be on a webserver connected to the internet
- To access the web-site we use client software such as a browser, which can interpret the content of the files fetched from the server
- The HTTP-protocol describes how the client and server should communicate





# Url (request)

- When we write an url in the address bar or click a link in the browser, the browser performs a HTTP-request to the server
- The server returns the content of the requested file along with some meta-data (http-response)
- The browser interprets the HTML in the response and renders it on the screen
- If the HTML contains references to images, stylesheets, javascript-files etc. these are fetched as well
- Url-structure:

protocol      Server-address      path

https://manning.com/index.html

Standard port: 80

URL	Status	Domain	Size	Remote IP	Timeline
GET index.html	200 OK	manning.com	95,3 KB	184.173.95.35:80	788ms
GET mainhome.css	200 OK	manning.com	12,0 KB	184.173.95.35:80	264ms
GET brand?form=cse-search-b	302 Found	google.com	260 B	216.58.209.100:80	36ms
GET navblog.js	200 OK	manning.com	4,2 KB	184.173.95.35:80	266ms
GET masthead.jpg	200 OK	manning.com	12,5 KB	184.173.95.35:80	392ms
GET top_button_ordering.jpg	200 OK	manning.com	1,2 KB	184.173.95.35:80	388ms
GET top_button_shoppingcart.j	200 OK	manning.com	1,6 KB	184.173.95.35:80	389ms
GET brand?form=cse-search-b	200 OK	cse.google.com	1,1 KB	216.58.209.110:80	240ms
GET google_custom_search_wa	200 OK	cse.google.com	2,0 KB	216.58.209.110:80	19ms
GET mailbox.gif	200 OK	manning.com	74 B	184.173.95.35:80	132ms
GET twitter-16x16.png	200 OK	manning.com	3,0 KB	184.173.95.35:80	131ms
GET facebookF.gif	200 OK	manning.com	1,0 KB	184.173.95.35:80	206ms
GET gupta2_3d.gif	200 OK	manning.com	5,8 KB	184.173.95.35:80	203ms
GET bdavis_3d.gif	200 OK	manning.com	6,7 KB	184.173.95.35:80	203ms
GET nicolette_3d.gif	200 OK	manning.com	7,7 KB	184.173.95.35:80	249ms

# HTML

- Describes the structure of a page/document
- The first version of HTML was only intended for online documents and not web-pages with complicated layouts
- The structure is a hierarchy of nested elements
  - An element is made of a start- and end-tag
  - Most elements have semantics
- Start-tags can have attributes

```
<!DOCTYPE html>
<html>
<head>
  <meta charset="utf-8">
  <title>A title here</title>
</head>
<body>
  <h1>Heading</h1>
  <p>
    some regular text <a href="http://dr.dk">with a link</a>
    and an image
    
  </p>
</body>
</html>
```

# Live Code –HTML

- Create a document with a few articles.
- Each article has:
  - A heading (<h1>)
  - body-text (<p>
    - With some links (<a href="LINK">)
  - An image ()

# Live Code Solution

## News of The World



### HPV vaccine cutting cervical cancer by nearly 90%

[bbc](https://www.bbc.com/news/health-59148620) says that Cancer Research UK described the findings as "historic", and said it showed the vaccine was saving lives.



### In a Supertall Tower, How Much Affordable Housing Is Enough?

[ny-times](https://www.nytimes.com/2021/10/29/realestate/5-world-trade-center-affordable.html) raises the question if the only residential building at the World Trade Center will be 25 percent affordable — a real accomplishment, supporters say. Others insist it should be 100 percent.

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>News of the World</title>
</head>
<body>
  <h1>News of The World</h1>
  <article>
    <h2>HPV vaccine cutting cervical cancer by nearly 90%</h2>
    <p><a href="https://www.bbc.com/news/health-59148620">bbc</a> says that Cancer
Research UK described the findings as "historic", and said it showed the vaccine was saving lives. </p>
    
  </article>

  <article>
    <h2>In a Supertall Tower, How Much Affordable Housing Is Enough?</h2>
    <p><a href="https://www.nytimes.com/2021/10/29/realestate/5-world-trade-center-affordable.html">ny-times</a>
writes that the only residential building at the World Trade Center will be 25 percent affordable —
a real accomplishment, supporters say. Others insist it should be 100 percent.</p>
      
<tr><td>cell 1</td><td>cell 2</td></tr>  
</table>
```

```
<p class="attention">  
Some text  
</p>
```

```
<h1 class="attention">Heading 1</h1>
```

cell 1 cell 2

Some text

**Heading 1**



# Id Selectors

- Can only be applied to one element in each HTML document
- Is applied using the id attribute
- Can coexist with the class attribute
- Id is more specific than class and therefore takes precedence if the same property is declared in both class- and id-rule

Id- selectors starts with "#"

```
<p class="attention" id="introduction">  
Some text  
</p>
```

*Some text*

# Context Specific Selectors

- descendant selector -> elements which matches the selector and is descendant of the outer selector
- child selector -> elements which matches the selector and is direct child of the outer selector

```
p .mini {  
  font-size: 10px;  
  background-color: red;  
}  
  
p > em {  
  color: white;  
  background-color: black;  
}
```

<p>  
 With <em>innovative</em>  
 <span> new <em>performance</em></span>  
 features,  
 <span><em class="mini">advanced</em></span>  
 layering and synchronization  
</p>

With **innovative** new performance features,  
**advanced** layering and synchronization

# Live code - CSS

- Apply styling to our web-page

## News of The World



### HPV vaccine cutting cervical cancer by nearly 90%

[bbc](#) says that Cancer Research UK described the findings as "historic", and said it showed the vaccine was saving lives.



### In a Supertall Tower, How Much Affordable Housing Is Enough?

[ny-times](#) raises the question if the only residential building at the World Trade Center will be 25 percent affordable — a real accomplishment, supporters say. Others insist it should be 100 percent.

## News of The World



### HPV vaccine cutting cervical cancer by nearly 90%

[bbc](#) says that Cancer Research UK described the findings as "historic", and said it showed the vaccine was saving lives.



### In a Supertall Tower, How Much Affordable Housing Is Enough?

[ny-times](#) raises the question if the only residential building at the World Trade Center will be 25 percent affordable — a real accomplishment, supporters say. Others insist it should be 100 percent.

# Live Code Solution

```
h1 {
  color: #fff;
  background-color: #000;
  text-align: center;
  line-height: 2em;
}

article {
  margin: auto;
  width: 800px;
  border-bottom: 3px dashed hotpink;
  padding-bottom: 20px;
}

article:after {
  content: "";
  clear: both;
  display: table;
}

article p {
  line-height: 1.4em;
  margin-bottom: 0px;
}

article img {
  float: left;
  width: 200px;
  margin-right: 30px;
}
```

## News of The World



### HPV vaccine cutting cervical cancer by nearly 90%

[bbc](#) says that Cancer Research UK described the findings as "historic", and said it showed the vaccine was saving lives.



### In a Supertall Tower, How Much Affordable Housing Is Enough?

[ny-times](#) raises the question if the only residential building at the World Trade Center will be 25 percent affordable — a real accomplishment, supporters say. Others insist it should be 100 percent.

# Web-Scraping Python

- We use 2 modules
- requests
  - Request a HTML-document using an url
- BeautifulSoup4
  - Parse the html-page into a set of python-objects we can query using CSS-selector syntax

```
> pip install requests BeautifulSoup4
```

# requests usage

- requests can be used to make HTTP-requests and fetch the response
  - Documentation: <https://docs.python-requests.org/en/latest/>
- If the encoding seems wrong use:
  - `response.encoding = response.apparent_encoding`

```
import requests

response = requests.get('https://dr.dk')
# response.encoding = response.apparent_encoding
print(response.text)
```

# beautifulsoup4 usage

- bs4 parses an HTML-string into a structure of python objects which can be queried using CSS-selectors (like the DOM for JS)
  - Documentation: <https://www.crummy.com/software/BeautifulSoup/bs4/doc/>
  - <https://www.crummy.com/software/BeautifulSoup/bs4/doc/#css-selectors>
- The `select()` method is used to get a list of tag-objects
  - The whole document and each tag can be queried using `select()`
- The `text` property (or method `get_text()`) gives us the text-content of the tag and nested tags
- We can access attributes of a Tag-object using the `get()` method

```
import bs4
```

```
html = 'some html...'
```

```
soup = bs4.BeautifulSoup(html, 'html.parser')
```

```
ps = soup.select('p')
```

```
for p in ps:
```

```
    print(p.text)
```

```
    print(p.get('id'))
```

# Find the CSS-path

- In the browser
  - right-click the element we are interested in and click “inspect”
  - We can now see the HTML-structure for the element clicked
  - Try if the type/class/id selector is enough to isolate the wanted text
    - if not, add more context

The image shows a browser window with a news article. The article title is "Caroline flyttede ind til udsigt over Aarhus, men nu vil kommunen opføre boliger i fire etager". Below the title is a photo of a building. To the right of the photo is a small text block that says "PENG" and "Mærsk vil kræve medarbejdere vaccineret: Ekspert og fagforeninger kalder det ulovligt".

Below the article is a row of three small images. The first image is a close-up of a man's face. The second image is a close-up of a woman's face. The third image is a close-up of a man's face.

On the right side of the image is a screenshot of the browser's developer tools. The "Elements" panel is open, showing the HTML structure of the page. The selected element is a `<span>` element with the class `dre-title-text`. The CSS path for this element is shown as `span.dre-title-text`. The "Styles" panel is also open, showing the default styles for the `span` element.

```
span.dre-title-text 446.63 x 144
```

```
<ul class="dre-grid-layout dre-grid-layout--xxs-c0-c0-c1-c2--sm-c0-c0-c0-c0-c1-c1-c0-c0-c0-c0-c2-c2 dre-grid-layout--spacing-small dre-grid-layout--list">  
  <li class="dre-grid-layout-list__item">  
    <div class="dre-article-teaser">  
      <div class="dre-teaser">  
        <div class="dre-teaser-image">  
          <div class="dre-article-teaser__media-label">  
            <div class="dre-teaser-content">  
              <div class="dre-article-teaser__text-box dre-article-teaser__text-box--xxs-x-small dre-article-teaser__text-box--xs-small dre-article-teaser__text-box--lg-medium">  
                <div class="dre-article-teaser__meta-wrapper">  
                  <a href="/nyheder/politik/kommunalvalg/caroline-flyttede-ind-til-udsigt-over-aarhus-men-nu-vil-kommunen" class="dre-teaser-title dre-teaser-title--xxs-small dre-teaser-title--xs-medium dre-teaser-title--md-large" aria-label="Caroline flyttede ind til udsigt over Aarhus, men nu vil kommunen opføre boliger i fire etager, fra sektionen "Kommunal- og regionalvalg">  
                    <span class="dre-teaser-title__text">  
                      <span class="dre-title-text">  
                        </span>  
                      </span>  
                    </a>  
                  </div>  
                </div>  
              </div>  
            </li>  
          </ul>
```



# Live Code – Web Scraping

- Fetch all article headings from dr.dk
- Fetch the title and year-span of the first 100 collections matching “Hasselager” from arkiv.dk
  - We need to paginate
  - We need to break when we have found the first 100 or if there are no more results

PROGRAMMING



# dr.dk solution

```
import requests
import bs4

def scrape():
    response = requests.get("https://dr.dk")
    soup = bs4.BeautifulSoup(response.text, 'html.parser')
    for h2 in soup.select('.dre-teaser-title'):
        print(h2.text)

if __name__ == "__main__":
    scrape()
```

# Arkiv.dk solution

```
import requests
import bs4

def scrape():
    page = 1
    count = 0
    while True:
        response = requests.get("https://arkiv.dk/soeg?searchstring=Hasselager&page=" + str(page))
        soup = bs4.BeautifulSoup(response.text, 'html.parser')
        rows = soup.select('#resultater tr')
        for row in rows:
            divs = row.select('div > div')
            year_span = divs[0].text.strip().replace(' ', '')
            title = divs[1].text.strip()
            print(year_span)
            print(title)
            print(100 * '-')
            count += 1
            if count == 100:
                break

        if len(rows) == 0 or count == 100: # no more collections
            break
        else:
            page += 1

if __name__ == "__main__":
    scrape()
```

# Exercises

- Extract all headlines from <https://tv2.dk>
- Extract all names and titles of researchers from <https://www.au.dk/om/presse/ekspertlister/corona>
- Extract all 116 titles from <http://comicscontainer.dk/comicbooks?dynamicSearch=ru&page=1>
  - Consider how you will determine when you reached the last page (last page + 1), so you don't get an infinite loop
  - The markup is not very well-structured so you probably need to make several sub-queries with bs4